

Some thoughts on security after ten years of qmail 1.0

Daniel J. Bernstein

Department of Mathematics, Statistics, and Computer Science (M/C 249)
University of Illinois at Chicago, Chicago, IL 60607-7045, USA
djb@cr.yp.to

ABSTRACT

The qmail software package is a widely used Internet-mail transfer agent that has been covered by a security guarantee since 1997. In this paper, the qmail author reviews the history and security-relevant architecture of qmail; articulates partitioning standards that qmail fails to meet; analyzes the engineering that has allowed qmail to survive this failure; and draws various conclusions regarding the future of secure programming.

Categories and Subject Descriptors

D.2.11 [Software Engineering]: Software Architectures—*bug elimination, code elimination*; D.4.6 [Operating Systems]: Security and Protection; H.4.3 [Information Systems Applications]: Communications Applications—*electronic mail*

General Terms

Security

Keywords

Eliminating bugs, eliminating code, eliminating trusted code

1. INTRODUCTION

1.1 The bug-of-the-month club

Every Internet service provider runs an MTA (a “Mail Transfer Agent”). The MTA receives mail from local users; delivers that mail to other sites through SMTP (the Internet’s “Simple Mail Transfer Protocol”); receives mail from other sites through SMTP; and delivers mail to local users.

I started writing an MTA, qmail, in 1995, because I was sick of the security holes in Eric Allman’s “Sendmail” software. Sendmail was by far the most popular MTA on the

*Date of this document: 2007.11.01. Permanent ID of this document: acd21e54d527dabf29afac0a2ae8ef41.

Internet at the time; see, e.g., [4]. Here’s what I wrote in the qmail documentation in December 1995:

Every few months CERT announces Yet Another Security Hole In Sendmail—something that lets local or even remote users take complete control of the machine. I’m sure there are many more holes waiting to be discovered; Sendmail’s design means that any minor bug in 41000 lines of code is a major security risk. Other popular mailers, such as Smail, and even mailing-list managers, such as Majordomo, seem just as bad.

Fourteen Sendmail security holes were announced in 1996 and 1997. I stopped counting after that, and eventually I stopped paying attention. Searches indicate that Sendmail’s most recent emergency security release was version 8.13.6 in March 2006; see [10] (“remote, unauthenticated attacker could execute arbitrary code with the privileges of the Sendmail process”).

After more than twenty years of Sendmail releases known to be remotely exploitable, is anyone willing to bet that the latest Sendmail releases are not remotely exploitable? The announcement rate of Sendmail security holes has slowed, but this fact doesn’t help the administrators whose systems have been broken into through Sendmail security holes.

1.2 The qmail release

I started serious code-writing for qmail in December 1995. I had just finished teaching a course on algebraic number theory and found myself with some spare time. The final kick to get something done was a promise I had made to a colleague: namely, that I would run a large mailing list for him. Sendmail didn’t offer me the easy list administration that I wanted, and it seemed to take forever to deliver a message (serially!) to a long list of recipients, never mind Sendmail’s reliability problems and security problems.

The 7 December 1995 version of qmail had 14903 words of code, as measured by

```
cat *.c *.h | cpp -fpreprocessed \
| sed 's/[a-zA-Z0-9][a-zA-Z0-9]*\//g' \
| tr -d '\012' | wc -c
```

(Other complexity metrics paint similar pictures.) The 21 December 1995 version had 36062 words. The 21 January 1996 version, qmail 0.70, had 74745 words. After watching this version run on my computer for a few days I released it to the public, starting the qmail beta test.

On 1 August 1996 I released qmail 0.90, 105044 words, ending the qmail beta test. On 20 February 1997 I released

qmail 1.00, 117685 words. On 15 June 1998 I released the current version, qmail 1.03, 124540 words. A slight derivative created by the community, netqmail 1.05, has 124911 words. I am aware of four bugs in the qmail 1.0 releases.

For comparison: Sendmail 8.7.5, released in March 1996, had 178375 words; Sendmail 8.8.5, released in January 1997, had 209955 words; Sendmail 8.9.0, released in May 1998, had 232188 words. The Sendmail release notes report hundreds of bugs in these releases. There are some user-visible feature differences between Sendmail and qmail, such as qmail's POP support, and Sendmail's UUCP support, and qmail's user-controlled mailing lists, and Sendmail's “remote root exploit” feature—just kidding!—but these don't explain the complexity gap; most of the code in each package is devoted to core MTA features needed at typical Internet sites.

Fingerprinting indicates that more than a million of the Internet's SMTP servers run either qmail 1.03 or netqmail 1.05. The third-party qmail.org site says

A number of large Internet sites are using qmail: USA.net's outgoing email, Address.com, Rediffmail.com, Colonize.com, Yahoo! mail, Network Solutions, Verio, MessageLabs (searching 100M emails/week for malware), listserv.acsu.buffalo.edu (a big listserv hub, using qmail since 1996), Ohio State (biggest US University), Yahoo! Groups, Listbot, USWest.net (Western US ISP), Telenordia, gmx.de (German ISP), NetZero (free ISP), Critical Path (email outsourcing service w/ 15M mailboxes), PayPal/Confinity, Hypermart.net, Casema, Pair Networks, Topica, MyNet.com.tr, FSmail.net, Mycom.com, and vuurwerk.nl.

Several authors have written qmail books: [7], [20], [14], [22]. Comprehensive statistics are hard to collect, but samples consistently indicate that qmail sends and receives a large fraction of all of the legitimate email on the Internet.

1.3 The qmail security guarantee

In March 1997, I took the unusual step of publicly offering \$500 to the first person to publish a verifiable security hole in the latest version of qmail: for example, a way for a user to exploit qmail to take over another account. My offer still stands. Nobody has found any security holes in qmail. I hereby increase the offer to \$1000.

Of course, “security hole in qmail” does not include problems outside of qmail: for example, NFS security problems, TCP/IP security problems, DNS security problems, bugs in scripts run from `.forward` files, and operating-system bugs generally. It's silly to blame a problem on qmail if the system was already vulnerable before qmail was installed!

It's not as silly to blame qmail for failing to encrypt and authenticate mail messages sent through the network; maybe cryptography should be handled by applications such as qmail, rather than by TCP/IP. But cryptography is outside the scope of the qmail security guarantee.

Denial-of-service attacks are also specifically disallowed: they are present in every MTA, widely documented, and very hard to fix without a massive overhaul of several major protocols. One could argue, and I would agree, that Internet mail desperately needs this overhaul; but that's a topic for another paper.

What the qmail security guarantee *does* say is that users

can't exploit qmail to steal or corrupt other users' data. If other programs met the same standard, and if our network links were cryptographically protected, then the only remaining security problems on the Internet would be denial-of-service attacks.

1.4 Contents of this paper

How was qmail engineered to achieve its unprecedented level of security? What did qmail do well from a security perspective, and what could it have done better? How can we build other software projects with enough confidence to issue comparable security guarantees?

My views of security have become increasingly ruthless over the years. I see a huge amount of money and effort being invested in security, and I have become convinced that most of that money and effort is being wasted. Most “security” efforts are designed to stop yesterday's attacks but fail completely to stop tomorrow's attacks and are of no use in building invulnerable software. These efforts are a distraction from work that *does* have long-term value.

In retrospect, some of qmail's “security” mechanisms were half-baked ideas that didn't actually accomplish anything and that could have been omitted with no loss of security. Other mechanisms have been responsible for qmail's successful security track record. My main goal in this paper is to explain how this difference could have been recognized in advance—how software-engineering techniques can be *measured* for their long-term security impact.

Section 2 articulates three specific directions of progress towards invulnerable software systems. The remaining sections of the paper discuss qmail's successes and failures in these three directions.

Much of what I say has been said many times before. (This isn't the first paper on software security; it isn't even the first paper on qmail security.) I apologize for not having taken the time to locate original sources.

2. HOW CAN WE MAKE PROGRESS?

A software *bug* means, by definition, a software feature violating the user's requirements. A software *security hole* means, by definition, a software feature violating the user's security requirements. Every security hole is therefore a bug.

(Advocates of formal specifications, security policies, and so on will correctly point out that making a complete list of user requirements is quite difficult—especially in gray areas that the user hasn't thought through. This complexity has an effect on the difficulty of eliminating security holes, as discussed below. But the difficulty would exist even without the complexity: today's software doesn't even meet the most basic security requirements that come to mind!)

Suppose that there is, on average, 1 bug in every N words of code—but that there are $10000N$ words of code inside the computer. The unhappy conclusion is that the computer has about 10000 bugs. Many of those bugs are, presumably, security holes.

How can we make progress towards having *no* security holes? How can we *measure* the progress? This section gives three answers. This section also discusses three ways that the community has distracted itself from making progress.

2.1 Answer 1: eliminating bugs

The first answer, and surely the most obvious answer, is to reduce the bug rate.

We can estimate the bug rate of a software-engineering process by carefully reviewing the resulting code and tracking the number of bugs found as a function of the amount of code reviewed. We can then compare the bug rates of different software-engineering processes. We can meta-engineer processes with lower bug rates.

(Sometimes subtle bugs slip past a code review. However, experience suggests that the overall bug rate, taking account of all the user requirements, is only slightly larger than the not-so-subtle-bug rate. “Given enough eyeballs, all bugs are shallow,” Eric Raymond commented in [17, Section 4: “Release Early, Release Often”].)

Bug elimination is one of the classic topics of software-engineering research. It is well known, for example, that one can drastically reduce the bug rate of a typical software-engineering process by adding coverage tests. It does not seem to be as well known that bug rates are affected by choices in earlier stages of the software-engineering process. See Section 3 of this paper for further discussion.

2.2 Answer 2: eliminating code

Consider again the computer system with $10000N$ words of code. Suppose that the bug rate is reduced far below 1 bug in every N words of code, but not nearly far enough to eliminate all the bugs in $10000N$ words of code. How can we make progress?

The second answer is to reduce the amount of code in the computer system. Software-engineering processes vary not only in the number of bugs in a given volume of code, but also in the volume of code used to provide the features that the user wants. Code elimination is another classic topic of software-engineering research: we can compare the code volumes produced by different software-engineering processes, and we can meta-engineer processes that do the job with lower volumes of code. See Section 4 of this paper for further discussion.

Rather than separately measuring the bugs-code ratio and the code-job ratio one could measure the product, the bugs-job ratio. However, the separate measures seem to highlight different techniques. Furthermore, the code-job ratio is of independent interest as a predictor of software-engineering time, and the bugs-code ratio is of independent interest as a predictor of debugging time.

2.3 Answer 3: eliminating trusted code

Suppose that, by producing a computer system with less code, and producing code with fewer bugs, we end up with a system having fewer than 1000 bugs. Presumably many of those bugs are still security holes. How can we make progress?

The third answer is to reduce the amount of *trusted* code in the computer system. We can architect computer systems to place most of the code into *untrusted* prisons. “Untrusted” means that code in these prisons—no matter what the code does, no matter how badly it behaves, no matter how many bugs it has—cannot violate the user’s security requirements. We can measure the amount of *trusted* code in our computer systems, and we can meta-engineer processes that produce systems with lower volumes of trusted code.

Of course, general techniques for reducing the amount of code in a system are also helpful for reducing the amount of trusted code. But additional techniques allow the amount of trusted code to be much smaller than the total amount

of code. There is a pleasant synergy between eliminating trusted code and eliminating bugs: we can afford relatively expensive techniques to eliminate the bugs in trusted code, simply because the volume of code is smaller.

Consider, for example, the portion of the `sendmail` code responsible for extracting an email address from the header of a mail message. According to [12], Mark Dowd discovered in 2003 that this code had a security hole:

Attackers may remotely exploit this vulnerability to gain “root” or superuser control of any vulnerable Sendmail server. . . . This vulnerability is especially dangerous because the exploit can be delivered within an email message and the attacker doesn’t need any specific knowledge of the target to launch a successful attack. . . . X-Force has demonstrated that this vulnerability is exploitable in real-world conditions on production Sendmail installations.

Suppose that the same address-extraction code is run under an interpreter enforcing two simple data-flow rules:

- the only way that the code can see the rest of the system is by reading this mail message;
- the only way that the code can affect the rest of the system is by printing one string determined by that mail message.

The code is then incapable of violating the user’s security requirements. An attacker who supplies a message that seizes complete control of the code can control the address printed by the code—but the attacker could have done this anyway without exploiting any bugs in the code. The attacker is incapable of doing anything else. See Section 5.2 for another example of the same idea.

(I am implicitly assuming that the user’s security requirements do not prohibit the creators of a mail message from controlling the address extracted from the message. But the requirements may be more complicated. Perhaps the user is creating a mail message by combining attachments from several sources; perhaps the user requires that each attachment be incapable of affecting other attachments, the header, etc. The interpreter then has to impose corresponding data-flow restrictions.)

By reducing the amount of trusted code far enough, and reducing the bug rate far enough, one can reasonably hope to produce a system where there are no bugs in the trusted code, and therefore no security holes. Presumably there are still bugs in the untrusted code, but those bugs cannot violate the user’s security requirements. Engineering a secure software system is easier than engineering a bug-free software system.

2.4 Distraction 1: chasing attackers

For many people, “security” consists of observing current attacks and changing something—anything!—to make those attacks fail. It is easy to understand the attraction of this reactive type of work: watching an attack fail, where earlier it would have succeeded, provides instant gratification for the defender.

Sometimes the changes fix the specific bugs exploited by the attacks. The Ubuntu distribution of the Linux operating system has issued more than 100 emergency security patches this year, fixing various bugs in a wide range of programs.

Sometimes the changes don't fix any bugs. "Firewalls" and "anti-virus systems" and "intrusion-detection systems" attempt to recognize attacks without patching the software targeted by the attacks.

Either way, the changes do nothing to fix the software-engineering deficiencies that led to the security holes being produced in the first place. If we define success as stopping yesterday's attacks, rather than as making progress towards stopping all possible attacks, then we shouldn't be surprised that our systems remain vulnerable to tomorrow's attacks.

2.5 Distraction 2: minimizing privilege

Many additional "security" efforts are applications of the "principle of least privilege." The principle is widely credited to Saltzer and Schroeder, who stated it as follows in [18]: "Every program and every user of the system should operate using the least set of privileges necessary to complete the job."

These "security" efforts work as follows. We observe that program P has no legitimate need to access operating-system resource R . We then use (and possibly extend) operating-system controls to prevent P from accessing R . We prevent an image-displaying program from sending data through the network; we prevent a DNS-lookup program from reading disk files; etc. See, for example, [3], [21], [11], [2], [15], [1], [16], and [23]. Section 5.1 discusses some qmail examples.

I have become convinced that this "principle of least privilege" is fundamentally wrong. Minimizing privilege might reduce the damage done by some security holes but almost never fixes the holes. Minimizing privilege is not the same as minimizing the amount of trusted code, does not have the same benefits as minimizing the amount of trusted code, and does not move us any closer to a secure computer system.

Consider, as an example, [11]'s confinement of Netscape's "DNS helper" program, preventing the program from accessing the local disk. This confinement did not prevent the *libresolv* bug in [8] from being a security hole in Netscape: an attacker could use the bug to seize control of the "DNS helper," modify all subsequent DNS data seen by Netscape, and steal the user's web connections. The situation before [11] was that bugs in the "DNS helper" had the power to violate the user's security requirements and therefore needed to be fixed; the situation after [11] was that bugs in the "DNS helper" had the power to violate the user's security requirements and therefore needed to be fixed.

The defining feature of untrusted code is that it *cannot* violate the user's security requirements. Turning a "DNS helper" into untrusted code is necessarily more invasive than merely imposing constraints upon the operating-system resources accessed by the program. The "DNS helper" handles data from many sources, and each source must be prevented from modifying other sources' data.

2.6 Distraction 3: speed, speed, speed

Programmers waste enormous amounts of time thinking about, or worrying about, the speed of noncritical parts of their programs, and these attempts at efficiency actually have a strong negative impact when debugging and maintenance are considered. We *should* forget about small efficiencies, say about 97% of the time; premature optimization is the root of all evil.

—Knuth in [13, page 268]

The most obvious effect of the pursuit of speed is that programmers put effort into low-level speedups, attempting to save time by tweaking small sections of code.

Programmers know when they're doing this. They see the increased programming time. They see the increased bug rate. They are generally quite happy to change their engineering process to skip almost all of this effort. Knuth's commentary seems to be aimed at novice programmers who don't understand how to use profiling tools.

Unfortunately, the pursuit of speed has other effects that are not as blatant and that are not as easy to fix.

Consider the address-extraction example in Section 2.3. Using an interpreter to impose simple data-flow restrictions on the address-extraction code would make bugs in the code irrelevant to security—a huge benefit. However, most programmers will say "Interpreted code is too slow!" and won't even try it.

Starting a new operating-system process could impose similar restrictions without an interpreter; see Section 5.2. However, most programmers will say "You can't possibly start a new process for each address extraction!" and won't even try it.

Anyone attempting to improve programming languages, program architectures, system architectures, etc. has to overcome a similar hurdle. Surely some programmer who tries (or considers) the improvement will encounter (or imagine) some slowdown in some context, and will then accuse the improvement of being "too slow"—a marketing disaster.

I don't like waiting for my computer. I really don't like waiting for someone else's computer. A large part of my research is devoted to improving system performance at various levels. (For example, my paper [6] is titled "Curve25519: new Diffie-Hellman speed records.") But I find security much more important than speed. We *need* invulnerable software systems, and we need them today, even if they are ten times slower than our current systems. Tomorrow we can start working on making them faster.

I predict that, once we all have invulnerable software systems, we'll see that security doesn't actually need much CPU time. The bulk of CPU time is consumed by a tiny fraction of our programs, and by a tiny fraction of the code within those programs; time spent on security verification will be unnoticeable outside these "hot spots." A typical hot spot spends millions of CPU cycles on data from a single source; modern compiler techniques, in some cases aided by proofs, will be able to hoist all the security verification out of the inner loops. The occasional hot spots with tricky security constraints, such as encryption of network packets, will be trusted bug-free code.

3. ELIMINATING BUGS

For many years I have been systematically identifying error-prone programming habits—by reviewing the literature, analyzing other people's mistakes, and analyzing my own mistakes—and redesigning my programming environment to eliminate those habits.

I had already made some progress in this direction when I started writing qmail, and I made further progress during the development of qmail 1.0. This doesn't mean I'm happy with the bug rate of the programming environment I used for qmail; my error rate has continued to drop in the last decade, and I see many aspects of the qmail programming environment as hopelessly obsolete.

Fortunately, my bug rate in the mid-1990s was low enough that—given the low volume of qmail code, as discussed in Section 4—there were only a few bugs in qmail 1.0. None of those bugs were security holes. This is the explanation for qmail’s exceptional security record. Note that qmail’s privilege minimization didn’t help at all; see Section 5.1 for further discussion.

This section discusses several examples of *anti-bug meta-engineering*: modifying programming languages, program structures, etc. to reduce bug rates.

3.1 Enforcing explicit data flow

The standard argument against global variables is that they can create hidden data flow, often surprising the programmer.

Consider, for example, the following bug fixed in qmail 0.74: “`newfield_datemake` would leave `newfield_date` alone if it was already initialized, even though `qmail-send` calls `newfield_datemake` anew for each bounce.”

I originally wrote the `newfield_datemake` function as part of the `qmail-inject` program, which sends exactly one outgoing message and needs to create exactly one Date field, conveniently stored in a global variable `newfield_date`. But then I reused the same function in the `qmail-send` program, which sends many outgoing messages (“bounce messages,” i.e., non-delivery reports) and needs to create a new Date field for each outgoing message. I had forgotten that `newfield_datemake` didn’t reset an existing `newfield_date`; this global variable ended up transmitting information—specifically, the obsolete Date field—from the previous message to the current message.

Hidden data flow is also at the heart of buffer-overflow bugs. The statement `x[i] = m` in C might appear at first glance to modify only the `x` variable but—if `i` is out of range—might actually modify any variable in the program, including return addresses and memory-allocation control structures. Similarly, reading `x[i]` might read any variable in the program.

Several aspects of qmail’s design make qmail’s internal data flow easier to see. For example, large portions of qmail run in separate UNIX processes. The processes are connected through pipelines, often through the filesystem, and occasionally through other communication mechanisms, but they do not have direct access to each other’s variables. Because each process has a relatively small state, it has relatively few opportunities for the programmer to screw up the data flow. At a lower level, I designed various array-access functions for which the indices were visibly in range, and I avoided functions for which this was hard to check.

Nowadays I am much more insistent on programming-language support for smaller-scale partitioning, sane bounds checking, automatic updates of “summary” variables (e.g., “the number of nonzero elements of this array”), etc. By “sane bounds checking” I don’t mean what people normally mean by “bounds checking,” namely raising an exception if an index is out of range; what I mean is automatic array extension on writes, and automatic zero-fill on reads. (Out of memory? See Section 4.2.) Doing the same work by hand is silly.

3.2 Simplifying integer semantics

Another surprise for the programmer is that `y` can be much smaller than `x` after `y = x + 1`. This happens if `x`

is the largest representable integer, typically $2^{31} - 1$; `y` will then be the smallest representable integer, typically -2^{31} .

The closest that qmail has come to a security hole was a potential overflow (pointed out by Georgi Guninski) of a 32-bit counter that I had failed to check. Fortunately, the counter’s growth was limited by the available memory, which in turn was limited by standard configuration; but the same 32-bit increment operation in another context could easily have caused a disastrous bug.

Similar comments apply to other integer operations. The operation semantics *usually* match the mathematical semantics that the programmer intends, but *occasionally* don’t. If I want to detect those occasions, I have to go to extra work to check for overflows. If I want to have sane mathematical semantics applied on those occasions—extending the integer range, and failing only if I run out of memory—I have to go to extra work to use a large-integer library.

Most programming environments are meta-engineered to make typical software easier to write. They should instead be meta-engineered to make *incorrect* software *harder* to write. An operation that is not exactly what I normally want should take *more* work to express than an operation that is exactly what I normally want. There are occasions when I really do want arithmetic modulo 2^{32} (or 2^{64}), but I am happy to do extra work on those occasions.

In some languages, `a + b` means exactly what it says: the sum of `a` and `b`. Often these languages are dismissed as being “too slow” for general use: an inner loop such as

```
for (i = 0; i < n; ++i) c[i] = a[i] + b[i];
```

suddenly involves `n` calls to an expensive high-precision-integer-arithmetic function such as `gmp_add()`. But it is not rocket science for a compiler to generate code that keeps track of the locations of large integers and that replaces the `gmp_add()` operations by machine operations in the typical case that all integers are small. Perhaps some slowdowns are more difficult to address, but—as in Section 2.6—we should *first* get the code right and *then* worry about its speed.

3.3 Avoiding parsing

I have discovered that there are two types of command interfaces in the world of computing: good interfaces and user interfaces.

The essence of user interfaces is *parsing*: converting an unstructured sequence of commands, in a format usually determined more by psychology than by solid engineering, into structured data.

When another programmer wants to talk to a user interface, he has to *quote*: convert his structured data into an unstructured sequence of commands that the parser will, he hopes, convert back into the original structured data.

This situation is a recipe for disaster. The parser often has bugs: it fails to handle some inputs according to the documented interface. The quoter often has bugs: it produces outputs that do not have the right meaning. Only on rare joyous occasions does it happen that the parser and the quoter both misinterpret the interface in the same way.

I made these comments in the original qmail documentation, along with two examples of how parsing and quoting were avoided in qmail’s extremely simple internal file structures and program-level interfaces. But I didn’t say anything about how parsing bugs and quoting bugs could be avoided when external constraints prohibit better interfaces.

Consider, for example, the following bug, fixed in qmail 0.74: “`qmail-inject` did not check whether `USER` needed quoting.” What `qmail-inject` was doing here was creating a `From` line showing the user’s name:

```
From: "D. J. Bernstein" <djb@cr.yp.to>
```

Normally the user’s account name, in this case `djb`, can be inserted verbatim before the `@`; but the name has to be quoted in a particular way if it contains unusual characters such as parentheses. My tests didn’t check specifically for unusual characters in `USER`, so they didn’t distinguish verbatim insertion from proper encoding.

Consider, as another example, the following format-string bug in the UNIX `logger` program:

```
Presumably logger on your system is doing
  syslog(pri,buf)
instead of the correct
  syslog(pri,"%s",buf)
... It is my guess that this does not constitute a
security hole in logger beyond a denial-of-service
attack, since an attacker would have an awfully
difficult time encoding reachable VM addresses
into printable ASCII characters, but without seeing
disassembled object code I can't be sure. Better
safe than sorry.
```

Tests that don’t put `%` into `buf` won’t notice the difference between `syslog(pri,buf)` and `syslog(pri,"%s",buf)`. (The quote is from [5], four years before format-string bugs were widely appreciated; see [19, Introduction].)

Verbatim copying of “normal” inputs—with ad-hoc quoting required for “abnormal” inputs—seems to be a universal feature of user interfaces. One way to catch the resulting bugs, in situations where the interface cannot be improved, is to systematically convert each quoting rule into another test.

3.4 Generalizing from errors to inputs

The following bug was fixed in qmail 0.90: “Failure to `stat .qmail-owner` was not an error.”

If Bob puts several addresses into `~bob/.qmail-buddies` then qmail will forward `bob-buddies` mail to those addresses. Delivery errors are sent back to the original sender by default, but Bob can direct errors to a different address by putting that address into `.qmail-buddies-owner`.

To check whether `.qmail-buddies-owner` exists, qmail uses the UNIX `stat()` function. If `stat()` says that the file exists, qmail directs errors to `bob-buddies-owner`. If `stat()` says that the file doesn’t exist, qmail uses the original sender address.

The problem is that `stat()` can fail temporarily. For example, `.qmail-buddies-owner` could be on a network filesystem that is momentarily unavailable. Before version 0.90, qmail would treat this error the same way as nonexistence—which would be wrong if the file actually exists. The only correct behavior is for qmail to give up and try again later.

My test suite didn’t cover the temporary-failure case. I had some tests for common cases, and I had some tricky tests going to extra effort to set up various error cases, but I didn’t have any tests manufacturing network-filesystem errors.

The volume of error-handling code can be drastically reduced, as discussed in Section 4.2, but computer systems

always have *some* error-handling code. How can we prevent bugs in code that is rarely exercised?

Testing would have been much easier if I had factored the code into (1) a purely functional protocol handler that could talk to anything, not just the filesystem, and (2) a simple wrapper that plugged `stat()` into the protocol handler. Feeding a comprehensive set of test cases to the protocol handler would then have been quite easy—and would have immediately caught this bug.

3.5 Can we really measure anti-bug progress?

There is an obvious difficulty in modelling user-interface research in general, and bug-elimination research in particular. The goal is to have users, in this case programmers, make as few mistakes as possible in achieving their desired effects. How do we model this situation—how do we model human psychology—except by experiment? How do we even *recognize* mistakes without a human’s help?

If someone can write a program to recognize a class of mistakes, great—we’ll incorporate that program into the user interface, eliminating those mistakes—but we still won’t be able to recognize the *remaining* mistakes. As a mathematician I’m bothered by this lack of formalization; I expect to be able to *define* a problem whether or not I can *solve* it.

Fortunately, research can and does proceed without models. We can observe humans to measure their programming-bug rates, even though we don’t know the algorithms that the humans are using. We can see that some software-engineering tools are bug-prone, and that others are not, without having any idea how to mathematically prove it.

4. ELIMINATING CODE

To this very day, idiot software managers measure ‘programmer productivity’ in terms of ‘lines of code produced,’ whereas the notion of ‘lines of code spent’ is much more appropriate.

—Dijkstra in [9, page EWD962–4]

This section discusses several examples of *code-volume-minimization meta-engineering*: changing programming languages, program structures, etc. to reduce code volume. As in Section 3, some of these examples were used in qmail and contributed to qmail’s low bug count, while other examples show that it is possible to do much better.

4.1 Identifying common functions

Here is a section of code from Sendmail (line 1924 of `util.c` in version 8.8.5):

```
if (dup2(fdv[1], 1) < 0)
{
    syserr("%s: cannot dup2 for stdout", argv[0]);
    _exit(EX_OSERR);
}
close(fdv[1]);
```

The `dup2()` function copies a file descriptor from one location to another; this `dup2()/close()` pattern moves a file descriptor from one location to another. There are several other instances inside Sendmail of essentially the same `dup2()/close()` pattern.

This particular pattern occurs only once in qmail, inside an `fd_move()` function called from a dozen other locations in the code:

```

int fd_move(int to,int from)
{
    if (to == from) return 0;
    if (fd_copy(to,from) == -1) return -1;
    close(from);
    return 0;
}

```

Most programmers would never bother to create such a small function. But several words of code are saved whenever one occurrence of the `dup2()/close()` pattern is replaced with one call to `fd_move()`; replacing a dozen occurrences saves considerably more code than were spent writing the function itself. (The function is also a natural target for tests.)

The same benefit scales to larger systems and to a huge variety of functions; `fd_move()` is just one example. In many cases an automated scan for common operation sequences can suggest helpful new functions, but even without automation I frequently find myself thinking “Haven’t I seen this before?” and extracting a new function out of existing code.

4.2 Automatically handling temporary errors

Consider the following excerpt from `qmail-local`:

```
if (!stralloc_cats(&dtline, "\n")) temp_nomem();
```

The `stralloc_cats` function changes a dynamically resized string variable `dtline` to contain the previous contents of `dtline` followed by a linefeed. Unfortunately, this concatenation can run out of memory. The `stralloc_cats` function then returns 0, and `qmail-local` exits via `temp_nomem()`, signalling the rest of the `qmail` system to try again later.

There are thousands of conditional branches in `qmail`. About half of them—I haven’t tried to count exactly—are doing nothing other than checking for temporary errors.

In many cases I built functions such as

```

void outs(s)
char *s;
{
    if (substdio_puts(&ss1,s) == -1) _exit(111);
}

```

to try an operation and exit the program upon temporary error. However, I didn’t—and don’t—like repeating the same work for each operation.

I could have pushed these tests into a relatively small number of bottom-level subroutines for memory allocation, disk reads, etc., exiting the program upon any temporary error. However, this strategy is unacceptable for long-running programs such as `qmail-send`. Those programs aren’t allowed to exit unless the system administrator asks them to!

I’ve noticed that many libraries and languages take the same strategy, rendering them similarly unacceptable for long-running programs. Maybe specialized types of programs should use specialized software-engineering environments, but I don’t view long-running programs as a specialized case; I question the wisdom of designing software-engineering environments that are unsuitable for building those programs.

Fortunately, programming languages can—and in some cases do—offer more powerful exception-handling facilities, aborting clearly defined subprograms and in some cases automatically handling error reports. In those languages I would be able to write

```
stralloc_cats(&dtline, "\n")
```

or simply

```
dtline += "\n"
```

without going to extra effort to check for errors. The reduced code volume would eliminate bugs; for example, the bug “if `ipme_init()` returned -1, `qmail-remote` would continue” (fixed in `qmail` 0.92) would not have had a chance to occur.

When I wrote `qmail` I rejected many languages as being much more painful than C for the end user to compile and use. I was inexplicably blind to the possibility of writing code in a better language and then using an automated translator to convert the code into C as a distribution language. Stroustrup’s `cfront`, the original compiler from C++ to C, is an inspirational example, although as far as I know it has never acquired exception-handling support.

4.3 Reusing network tools

UNIX has a general-purpose tool, `inetd`, that listens for network connections. When a connection is made, `inetd` runs another program to handle the connection. For example, `inetd` can run `qmail-smtpd` to handle an incoming SMTP connection. The `qmail-smtpd` program doesn’t have to worry about networking, multitasking, etc.; it receives SMTP commands from one client on its standard input, and sends the responses to its standard output.

Sendmail includes its own code to listen for network connections. The code is more complicated than `inetd`, in large part because it monitors the system’s load average and reduces service when there is heavy competition for the CPU.

Why does Sendmail not want to handle mail when the CPU is busy? The basic problem is that, as soon as Sendmail accepts a new message, it immediately goes to a lot of effort to figure out where the message should be delivered and to try delivering the message. If many messages show up at the same time then Sendmail tries to deliver all of them at the same time—usually running out of memory and failing at most of the deliveries.

Sendmail tries to recognize this situation by checking the load average. If the CPU is busy, Sendmail inserts new messages into a queue of not-yet-delivered messages. Sendmail has a background delivery mechanism that periodically runs through the queue, trying to deliver each message in turn. The background delivery mechanism never overloads the computer; if many messages arrive in the queue at once, the messages aren’t even noticed until the next queue run, and are then handled with highly limited parallelism.

The queue-run interval “is typically set to between fifteen minutes and one hour,” the documentation says; “RFC 1123 section 5.3.1.1 recommends that this be at least 30 minutes.” System administrators who set a very short queue-run interval, for example 30 seconds, find Sendmail trying each queued message thousands of times a day.

Why does Sendmail have a foreground delivery mechanism? Why does it not put *all* incoming messages into the queue? The answer is that—even when the CPU is *not* busy—queued messages are not delivered immediately. Putting all incoming messages into the queue would mean waiting for the next queue run before delivery; presumably users would sometimes see the delays and complain.

With `qmail`, a small amount of extra code notifies the background delivery mechanism, `qmail-send`, when a message is placed into the queue. The `qmail-send` program

instantly tries delivering the message (if it is not busy), and tries again later if necessary, on a reasonable schedule. Consequently, the motivation for a foreground delivery mechanism disappears; qmail has no foreground delivery mechanism. Furthermore, the motivation for checking the load average disappears; qmail is happy to run from `inetd`.

If I *did* want to check the load average, I would do so in a general-purpose tool like `inetd`, rather than repeating the code in each application.

4.4 Reusing access controls

I started using UNIX, specifically Ultrix, twenty years ago. I remember setting up my `.forward` to run a program that created a file in `/tmp`. I remember inspecting thousands of the resulting files and noticing in amazement that Sendmail had occasionally run the program under a uid other than mine.

Sendmail handles a user's `.forward` as follows. It first checks whether the user is allowed to read `.forward`—maybe the user has set up `.forward` as a symbolic link to a secret file owned by another user. It then extracts delivery instructions from `.forward`, and makes a note of them (possibly in a queue file to be handled later), along with a note of the user responsible for those instructions—in particular, the user who specified a program to run. This is a considerable chunk of code (for example, all of `safefile.c`, plus several scattered segments of code copying the notes around), and it has contained quite a few bugs.

Of course, the operating system already has its own code to check whether a user is allowed to read a file, and its own code to keep track of users. Why write the same code again?

When qmail wants to deliver a message to a user, it simply starts a delivery program, `qmail-local`, under the right uid. When `qmail-local` reads the user's delivery instructions, the operating system automatically checks whether the user is allowed to read the instructions. When `qmail-local` runs a program specified by the user, the operating system automatically assigns the right uid to that program.

I paid a small price in CPU time for this code reuse: qmail starts an extra process for each delivery. But I also avoided all the extra system calls that Sendmail uses to check permissions. Anyway, until someone shows me a real-world mail-delivery computer bottlenecked by qmail's `fork()` time, I'm certainly not going to spend extra code to reduce that time.

4.5 Reusing the filesystem

Suppose that the National Security Agency's SMTP server receives mail for the `efd-friends@nsa.gov` mailing list. How does the MTA find out that it is supposed to accept mail for `nsa.gov`? How does the MTA find the delivery instructions for `efd-friends`?

Evidently the MTA will look up the names `nsa.gov` and `efd-friends` in a database. Maybe the database isn't called a "database"; maybe it's called an "associative array" or something else; whatever it's called, it is capable of giving back information stored under names such as `efd-friends` by the system administrator and the mailing-list manager.

With Sendmail, names such as `nsa.gov` and various other configuration options are listed inside a "configuration file" having a fairly complicated format. Mailing lists such as `efd-friends` are listed inside an "aliases file." Looking up `nsa.gov` and `efd-friends` inside these files requires a considerable amount of parsing code.

Of course, the operating system already has code to store chunks of data under specified names and to retrieve the chunks later. The chunks are called "files"; the names are called "filenames"; the code is called "the filesystem." Why write the same code again?

With qmail, the delivery instructions for `efd-friends` are the contents of a file named `.qmail-efd-friends`. Finding or modifying those instructions is a simple matter of opening that file. Users already have tools for creating and managing files; there is no need for qmail to reinvent those tools.

I should have, similarly, put the `nsa.gov` configuration into `/var/qmail/control/domains/nsa.gov`, producing the same simplicity of code. I instead did something requiring slightly more complicated code: `nsa.gov` is a line in a file rather than a file in a directory. I was worried about efficiency: most UNIX filesystems use naive linear-time algorithms to access directories, and I didn't want qmail to slow down on computers handling thousands of domains. Most UNIX filesystems also consume something on the scale of a kilobyte to store a tiny file.

In retrospect, it was stupid of me to spend code—not just this file-parsing code, but also code to distribute message files across directories—dealing with a purely hypothetical performance problem that I had not measured as a bottleneck. Furthermore, to the extent that measurements indicated a bottleneck (as they eventually did for the message files on busy sites), I should have addressed that problem at its source, fixing the filesystem rather than complicating every program that uses the filesystem.

5. ELIMINATING TRUSTED CODE

5.1 Accurately measuring the TCB

"Even if all of these programs are completely compromised, so that an intruder has control over the `qmaild`, `qmails`, and `qmailr` accounts and the mail queue, he still can't take over your system," I wrote in the qmail documentation. "None of the other programs trust the results from these five." I continued in the same vein for a while, talking about privilege minimization as something helpful for security.

Let's think about this for a minute. Suppose that there *is* a bug in `qmail-remote` allowing an attacker to take control of the `qmailr` account. The attacker can then steal and corrupt the system's outgoing mail, even if network connections are completely protected by strong cryptography. This is a security disaster, and it needs to be fixed. The only way that qmail avoids this disaster is by avoiding bugs.

Similarly, I shouldn't have highlighted the small amount of code in qmail capable of affecting files owned by `root`. Almost all of the code in qmail is capable of affecting files owned by normal users—either on disk or in transit through the mail system—and is therefore in a position to violate the users' security requirements. The only way that this code avoids security holes is by avoiding bugs.

Programmers writing word processors and music players generally don't worry about security. But users expect those programs to be able to handle files received by email or downloaded from the web. Some of those files are prepared by attackers. Often the programs have bugs that can be exploited by the attackers.

When I taught a "UNIX security holes" course in 2004, I asked the students to find new security holes for their

homework. I ended up disclosing 44 security holes found by the students. Most of those were in programs that are usually—incorrectly—viewed as being outside the system’s trusted code base.

For example, Ariel Berkman discovered a buffer overflow in xine-lib, a movie-playing library. Users were at risk whenever they played movies downloaded from the web: if the movies were supplied by an attacker then the attacker could read and modify the users’ files, watch the programs that the users were running, etc.

“Secure” operating systems and “virtual machines” say that their security is enforced by a small base of trusted code. Unfortunately, a closer look shows that this “security” becomes meaningless as soon as one program handles data from more than one source. The operating system does nothing to stop one message inside a mail-reading program from stealing and corrupting other messages handled by the same program, for example, or to stop one web page inside a browser from stealing and corrupting other web pages handled by the same browser. The mail-reading code and browsing code are, in fact, part of the trusted code base: bugs in that code are in a position to violate the user’s security requirements.

5.2 Isolating single-source transformations

The `jpegtopnm` program reads a JPEG file, a compressed image, as input. It uncompresses the image, produces a bitmap as output, and exits. Right now this program is trusted: its bugs can compromise security. Let’s see how we can fix that.

Imagine running the `jpegtopnm` program in an “extreme sandbox” that doesn’t let the program do anything other than read the JPEG file from standard input, write the bitmap to standard output, and allocate a limited amount of memory. Existing UNIX tools make this sandbox tolerably easy for `root` to create:

- Prohibit new files, new sockets, etc., by setting the current and maximum `RLIMIT_NOFILE` limits to 0.
- Prohibit filesystem access: `chdir` and `chroot` to an empty directory.
- Choose a uid dedicated to this process ID. This can be as simple as adding the process ID to a base uid, as long as other system-administration tools stay away from the same uid range.
- Ensure that nothing is running under the uid: `fork` a child to run `setuid(targetuid)`, `kill(-1, SIGKILL)`, and `_exit(0)`, and then check that the child exited normally.
- Prohibit `kill()`, `ptrace()`, etc., by setting gid and uid to the target uid.
- Prohibit `fork()`, by setting the current and maximum `RLIMIT_NPROC` limits to 0.
- Set the desired limits on memory allocation and other resource allocation.
- Run the rest of the program.

At this point, unless there are severe operating-system bugs, the program has no communication channels other than its initial file descriptors.

Suppose an attacker supplies a rogue JPEG file that exploits a bug in `jpegtopnm` and succeeds in seizing complete control of the program. The attacker can then generate any output bitmap he wants—but this is what he could have

done without exploiting any bugs. The attacker cannot do anything else. At this point `jpegtopnm` is no longer trusted code; bugs in `jpegtopnm` are no longer capable of violating the user’s security requirements.

As in Section 2.3, I am implicitly assuming that the user’s security requirements allow anyone who influences the JPEG file to control the resulting bitmap. This is a constraint on the security requirements, but I think it is a reasonable one; I have never heard anyone ask for pieces of a JPEG to be protected from each other. More importantly, I see no reason to believe that variations in the security requirements cannot be accommodated by analogous variations in the partitioning of code into sandboxes.

I am also assuming that the CPU does not leak secret information from one process to another. This assumption is debatable. Perhaps access to CPU instructions needs to be limited, for example with the interpreter discussed in Section 2.3. See Section 2.6 for comments on performance.

Ariel Berkman, following my suggestions, has reengineered the standard UNIX `xloadimage` picture-viewing tool, modularizing essentially the entire program into a series of filters such as `jpegtopnm`. Each filter is easily imprisoned by the techniques described above, leaving a much smaller amount of trusted code.

5.3 Delaying multiple-source merges

Consider a mail reader displaying messages from a user’s mailbox. Single-source transformations—for example, uncompressing an attached JPEG file—can be imprisoned as described in Section 5.2. But what about transformations that combine data from multiple sources?

A displayed list of message subjects combines information from several messages. Each message is allowed to control its own entry in the list, but it is not allowed to affect other entries. Even if the list is created correctly in the first place, bugs in subsequent transformations of the list can violate the required separation between list entries.

The situation is quite different if the subjects of separate messages are kept in separate locations, transformed independently, and *then* merged. The transformations are now single-source transformations. Delaying the merge reduces the volume of trusted code.

Stopping cross-site scripting, for example, currently means paying careful attention in every piece of code that merges data from multiple sources into one file, and in every piece of code helping transform that file into a web page. If data from different sources were instead kept separate, and merged only at the last moment by the browser, then care would be required only in the final merging code.

5.4 Do we really need a small TCB?

I failed to place any of the qmail code into untrusted prisons. Bugs anywhere in the code could have been security holes. The way that qmail survived this failure was by having very few bugs, as discussed in Sections 3 and 4.

Perhaps continued reductions of overall code volume and of bug rate will allow the same survival to scale to much larger systems. I’ve heard reports of systems that are believed to be bug-free despite having half a million lines of code. There are many more lines of code in my laptop computer—many more lines of code in a position to violate my security requirements—but it’s *conceivable* that the world could eliminate all the bugs in that code.

However, Section 5.2 shows that a large chunk of code can be eliminated from the TCB at very low cost—surely lower cost, and higher confidence, than eliminating bugs from the code. I’m optimistic about the scalability of this example; I don’t know exactly how small the ultimate TCB will be, but I’m looking forward to finding out. Of course, we still need to eliminate bugs from the code that remains!

6. REFERENCES

- [1] Anurag Acharya, Mandar Raje, *MAPbox: using parameterized behavior classes to confine untrusted applications*, 9th USENIX Security Symposium (2000). URL: <http://www.usenix.org/publications/library/proceedings/sec2000/acharya.html>. Citations in this document: §2.5.
- [2] Vinod Anupam, Alain Mayer, *Security of web browser scripting languages: vulnerabilities, attacks, and remedies*, 7th USENIX Security Symposium (1998). URL: <http://www.usenix.org/publications/library/proceedings/sec98/anupam.html>. Citations in this document: §2.5.
- [3] M. Lee Badger, Daniel F. Sterne, David L. Sherman, Kenneth M. Walker, Sheila A. Haghight, *A domain and type enforcement UNIX prototype*, 5th USENIX Security Symposium (1995). URL: <http://www.usenix.org/events/security95/badger.html>. Citations in this document: §2.5.
- [4] Daniel J. Bernstein, *Internet host SMTP server survey*, posting to comp.security.unix, comp.mail.misc, comp.mail.sendmail (1996). URL: <http://cr.yp.to/surveys/smtpsoftware.txt>. Citations in this document: §1.1.
- [5] Daniel J. Bernstein, *Re: Logging question* (1996). URL: <http://www.ornl.gov/lists/mailing-lists/qmail/1996/12/msg00314.html>. Citations in this document: §3.3.
- [6] Daniel J. Bernstein, *Curve25519: new Diffie-Hellman speed records*, in [24] (2006), 207–228. URL: cr.yp.to/papers.html#curve25519. Citations in this document: §2.6.
- [7] Richard Blum, *Running qmail*, Sams Publishing, 2000. ISBN 978-0672319457. Citations in this document: §1.2.
- [8] Computer Emergency Response Team, *CERT advisory CA-2002-19: buffer overflows in multiple DNS resolver libraries* (2002). URL: <http://www.cert.org/advisories/CA-2002-19.html>. Citations in this document: §2.5.
- [9] Edsger W. Dijkstra, *Introducing a course on mathematical methodology*, EWD962 (1986). URL: <http://www.cs.utexas.edu/users/EWD/ewd09xx/EWD962.PDF>. Citations in this document: §4.
- [10] Jeff Gennari, *Vulnerability Note VU#834865: Sendmail contains a race condition* (2006). URL: <http://www.kb.cert.org/vuls/id/834865>. Citations in this document: §1.1.
- [11] Ian Goldberg, David Wagner, Randi Thomas, Eric Brewer, *A secure environment for untrusted helper applications (confining the wily hacker)*, 6th USENIX Security Symposium (1996). URL: <http://www.usenix.org/publications/library/proceedings/sec96/goldberg.html>. Citations in this document:
- §2.5, §2.5, §2.5, §2.5.
- [12] Internet Security Systems, *Remote Sendmail header processing vulnerability* (2003). URL: <http://www.iss.net/issEn/delivery/xforce/alertdetail.jsp?oid=21950>. Citations in this document: §2.3.
- [13] Donald Knuth, *Structured programming with go to statements*, Computing Surveys **6** (1974), 261–301. Citations in this document: §2.6.
- [14] John R. Levine, *qmail*, O’Reilly, 2004. ISBN 978-1565926288. Citations in this document: §1.2.
- [15] Nimisha V. Mehta, Karen R. Sollins, *Expanding and extending the security features of Java*, 7th USENIX Security Symposium (1998). URL: <http://www.usenix.org/publications/library/proceedings/sec98/mehta.html>. Citations in this document: §2.5.
- [16] David S. Peterson, Matt Bishop, Raju Pandey, *A flexible containment mechanism for executing untrusted code*, 11th USENIX Security Symposium (2002). URL: <http://www.usenix.org/events/sec02/peterson.html>. Citations in this document: §2.5.
- [17] Eric Raymond, *The cathedral and the bazaar* (1997). URL: <http://www.catb.org/~esr/writings/cathedral-bazaar/cathedral-bazaar/>. Citations in this document: §2.1.
- [18] Jerry H. Saltzer, Mike D. Schroeder, *The protection of information in computer systems*, Proceedings of the IEEE **63** (1975), 1278–1308. Citations in this document: §2.5.
- [19] Umesh Shankar, Kunal Talwar, Jeffrey S. Foster, David Wagner, *Detecting format string vulnerabilities with type qualifiers*, 10th USENIX Security Symposium (2001). URL: <http://www.usenix.org/publications/library/proceedings/sec01/shankar.html>. Citations in this document: §3.3.
- [20] Dave Sill, *The qmail handbook*, Apress, 2002. ISBN 978-1893115408. Citations in this document: §1.2.
- [21] Kenneth M. Walker, Daniel F. Sterne, M. Lee Badger, Michael J. Petkac, David L. Sherman, Karen A. Oostendorp, *Confining root programs with domain and type enforcement*, 6th USENIX Security Symposium (1996). URL: <http://www.usenix.org/events/sec96/walker.html>. Citations in this document: §2.5.
- [22] Kyle Wheeler, *Qmail quickstarter: install, set up and run your own email server*, Packt Publishing, 2007. ISBN 978-1847191151. Citations in this document: §1.2.
- [23] Chris Wright, Crispin Cowan, Stephen Smalley, James Morris, Greg Kroah-Hartman, *Linux security modules: general security support for the Linux kernel*, 11th USENIX Security Symposium (2002). URL: <http://www.usenix.org/events/sec02/wright.html>. Citations in this document: §2.5.
- [24] Moti Yung, Yevgeniy Dodis, Aggelos Kiayias, Tal Malkin (editors), *9th international conference on theory and practice in public-key cryptography, New York, NY, USA, April 24–26, 2006, proceedings*, Lecture Notes in Computer Science, 3958, Springer, Berlin, 2006. ISBN 978-3-540-33851-2. See [6].